Complexity of the Gale String Problem for Equilibrium Computation in Games

Marta Maria Casetti

London, April 2015

Declaration

I certify that chapter 2 of this thesis I have presented for examination for the MPhil degree of the London School of Economics and Political Science is based on joint work with Julian Merschen and Bernhard von Stengel, published in [3].

The copyright of this thesis rests with the author. Quotation from it is permitted, provided that full acknowledgement is made. This thesis may not be reproduced without my prior written consent.

I warrant that this authorisation does not, to the best of my belief, infringe the rights of any third party.

Abstract

This thesis presents a report on original research, published as joint work with Merschen and von Stengel in *Electronic Notes in Discrete Mathematics* [3]. Our result shows a polynomial time algorithm to solve two problems related to labeled Gale strings, a combinatorial structure consisting a string of labels and a bitstring satisfying certain conditions introduced by Gale in [7].

Gale strings can be used in the representation of a particular class of games that Savani and von Stengel [17] used as an example of exponential running time for the classical Lemke-Howson algorithm to find a Nash equilibrium of a bimatrix game [9]. It was conjectured that solving these games via the Lemke-Howson algorithm was complete in the class **PPAD** (Proof by Parity Argument, Directed version). A major motivation for the definition of this class by Papadimitriou [16] was, in turn, to capture the pivoting technique of many results related to the Nash equilibrium, including the Lemke-Howson algorithm.

Our result, on the contrary, sets apart this class of games as a case for which there is a polynomial-time algorithm to find a Nash equilibrium. Since Daskalakis, Goldberg and Papaditrimiou [5] and Chen and Deng [4] proved the **PPAD**-completeness of finding a Nash equilibrium in general normal-form games, we have a special class of games, unless **PPAD** = **P**.

Our proof exploits two results. The first one is the representation of the Nash equilibria of these games as Gale strings, as seen in Savani and von Stengel [17]. The second one is the polynomial-time solvability of the problem of finding a perfect matching in a graph, proven by Edmonds [6].

Merschen [12] and Végh and von Stengel [20] expanded our technique to prove further interesting results.

Contents

List of Figures

Introduction

The topic of this thesis is a problem in the field of *algorithmic game theory*, that is, the study of game-theoretic problems from the point of view of computer science. In particular, we focus on the computational complexity of a particular class of games. These

General refs for comp compl [15]

General refs for geometry [23]

summary of chapt 2:

We begin this section with the definition of almost complete labeling; we then move on to the classic version of the Lemke-Howson algorithm for the problem ANOTHER COMPLETELY LABELED VERTEX, as given in the beautiful exposition by Shapley [19], and its dual version for ANOTHER COMPLETELY LABELED FACET. Finally, we present the Lemke-Howson for Gale algorithm. In the next session we will tackle the issue of the computational complexity of these algorithms: ANOTHER COMPLETELY LABELED FACET and ANOTHER COMPLETELY LABELED VERTEX are **PPA**, NASH is **PPAD**, as first shown in Papadimitriou [16]; furhermore, as shown by Morris [13] and by Savani and von Stengel [17], there are cases of exponential running time. This had led us to conjecture that these problems could be exploited for a proof of **PPAD** completeness, also considering that finding a completely labeled facet (or vertex, or the existence of a Nash equilibrium) is **NP** in the case of a general labeled polytope, as proven by von Stengel [22] . In the last section we will finally present our original result, that goes in the opposite direction:

check proof, citation

original? main?

the problem ANOTHER GALEcan be solved in polynomial time, that is, it is a problem in **TFP**. Unit vector games with dual cyclic best response polytope present therefore a case apart, as expected, but not because they are harder than others, but because they are easier.

Chapter 1

Polytopes, Games, Complexity, Labels and Gale Strings

New subdivision:

preliminaries: until (not including) TFNP;

the classes PPA and PPAD (add a bit more on PPA);

then labels and gale-def as before.

1.1 Preliminary definitions

We follow the notation $[n] = \{i \in \mathbb{N} \mid 1 \leq i \leq n\}$. We denote the transpose of a matrix A as $A^\top$. We consider vectors $u, v \in \mathbb{R}^d$ as column vectors, so $u^\top v$ is their scalar product. A vector in $\mathbb{R}^d$ for which all components are 0's will be denoted as $\mathbf{0}$; a vectors for which all components are 1's will be denoted as $\mathbf{1}$; the *i-th unit vector*, for which all the components are 0 except for the i-th component equal to 1, will be denoted as e_i. An inequality of the form $u \geq v$, is intended to hold for every component.

An *affine combination* of points in an Euclidean space $\{z_1, \ldots, z_m\} \subset^n$ is $\sum_{i=1}^m \lambda_i z_i$ where $\lambda_i \in \mathbb{R}$ such that $\sum_{i=1}^m \lambda_i = 1$; the points $z_1, \ldots, z_m$ are *affinely independent* if none of them is an affine combination of the others.

The *convex hull* of the points $z_1, \ldots, z_m$ is an affine with $\lambda_i \geq 0$ for all $i \in [n]$; that is $\mathrm{conv}\{z_1, \ldots, z_m\} = \{\sum_i \lambda_i z_i \mid \lambda_i \geq 0, \; \sum_i \lambda_i = 1$ A set of points $Z = \{z_1, \ldots, z_m\}$ is *convex* if $Z = \mathrm{conv}(Z)$; a convex set has *dimension d* if it has exactly $d+1$ affinely independent points; in this case we say that Z is a *d-simplex*. The *standard d-simplex* is denoted $\Delta_d = \mathrm{conv}\{e_1, \ldots, e_{d+1}\}$.

In general, a *polytope* is the convex hull of a finite set of points (not necessarily affinely independent) $\{z_1, \ldots, z_m\} \subset \mathbb{R}^n$; the *dimension* of the polytope is the dimension of its convex hull. A polyhedron is the intersection of finitely many closed halfspaces $\{x \in \mathbb{R}^d \mid a^T x \leq a_0\}$; a bounded polyhedron is, again, a polytope.

A *vertex* of a d-dimensional polytope $P = \mathrm{conv}(Z)$ is a point $z \in Z$ such that $\mathrm{conv}(Z \setminus \{z\}) \neq P$; an *edge* of P is a 1-dimensional line segment that has two vertices as endpoints. A *facet* of P is the convex hull of a set of d vertices $F = \{z_1, \ldots, z_d\}$ that lie on a hyperplane of the form $\{x \in \mathbb{R}^d \mid a^T x = a_0\}$ so that $a^T u < a_0$ for all other vertices u of P; the vector a (unique up to a scalar multiple) is called the *normal vector* of the facet.

A d-dimensional *simplicial polytope P* is the convex hull of a set of at least $d+1$ points $v \in \mathbb{R}^d$ such that no $d+1$ of them are on a common hyperplane; this is equivalent to requiring that every facet of P is a d-simplex. A d-dimensional polytope P is *simple* if every point of P lies on at most than d facets; the points that lie on exactly d facets will be the vertices.

The *polar* of the polytope $P = \{x \in \mathbb{R}^d \mid x^\top c_i \leq 1, \; i \in [k]\}$ with $c_i \in \mathbb{R}^d$ is $P^\Delta = \mathrm{conv}\{c_i, i \in [k]\}$.

1.2 Normal Form Games and Nash Equilibria

We now give the game-theoretic background that will be used in this thesis. A *game*, as first defined by von Neumann in [21], is a model of strategic interaction. A *finite normal form game* is $\Gamma = (P, S = \times_{p \in P} S_p, u = \times_{p \in P} u^p)$ where both the set of *players* P and the sets of *pure strategies* S_p (and there-

fore the set of *pure strategy profiles* S) are finite. We will use the notation $S_{-p} = \times_{q \neq p} S_p$. The purpose of each player $p \in P$ is to maximize her *payoff function* $u^p : S \to \mathbb{R}$. In the following pages, by "game" we will always mean "finite normal form game." If there are only two players, we will refer to player 1 using feminine pronouns and to player 2 using masculine ones; such games are called *bimatrix games* since they can be characterized by the $m \times n$ payoff matrices A and B, where a_{ij} and b_{ij} are the payoffs of respectively player 1 and of player 2 when the former plays her ith pure strategy and the latter plays his jth pure strategy. A bimatrix game is *zero-sum* if $B = -A$, and *symmetric* if $B = A^\top$.

A *mixed strategy* of player p is a probability distribution on S_P; it can be described as a point $x = (x_1^p, \ldots, x_{|S_p|}^p)$ on the $(|S_p| - 1)$-dimensional *mixed strategy simplex* $\Delta_p = \{x \in \mathbb{R}^{|S_p|} \mid x \geq \mathbf{0}, \ \mathbf{1}^\top x = 1\}$; the set of *mixed strategy profiles* will be the simplicial polytope $\Delta = \times_{p \in P} \Delta_p$. We extend the payoff functions to $u^p : \Delta \to \mathbb{R}$ by linearity.

A *Nash equilibrium* of a game is a strategy profile in which each player cannot improve her expected payoff by unilaterally changing her strategy; such a strategy is called a *best response*. Note that applying an affine transformation to all the payoffs does not change the Nash equilibria of the game.

Proposition 1.1. *A mixed strategy $x \in \Delta_p$ is a best response against some mixed strategy profile $y \in \Delta_{-p}$ of the other players if and only if every pure strategy $s_i \in S_p$ chosen with positive probability in x is a best response to y.*

The existence of a Nash equilibrium is guaranteed by the fundamental theorem by Nash ([14]). Note that there might be more than one equilibrium.

Theorem 1. (Nash [14]) *Every finite game in normal form has a Nash equilibrium.*

We give three classic examples of games: matching pennies, the prisoners' dilemma and a coordination game.

Example 1.1. Figure 1.1 shows the payoffs of the non-symmetric zero-sum game *matching pennies*, with payoff matrices $A = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$ and $B = -A$.

	s_1^2	s_2^2
s_1^1	-1 1	0 0
s_2^1	0 0	-1 1

Figure 1.1   Matching pennies.

At the unique equilibrium of the game, each player follows the uniform distribution over their strategies.

In the symmetric non zero-sum *prisoners' dilemma* of figure 1.2, each player must decide whether to "help" the other one or to "betray" them. If both players help each other, they will get a small reward; if both betray, they will pay a small penalty; if one betrays and the other cooperate the former will get a large reward and the latter will pay a large penalty.

	betray	help
betray	1 1	0 3
help	3 0	2 2

Figure 1.2   The prisoners' dilemma.

The only equilibrium is the profile in which both players betray. Assume that player 1 helps: then she must switch to betrayal, since she would get 10 instead of 5 if player 2 helps and -2 instead of -8 if player 2 betrays. The same applies to player 2, so both players will betray. The payoff matrices are $A = \begin{pmatrix} 1 & 3 \\ 0 & 2 \end{pmatrix}$ and $B = A^\top$.

Finally, in figure 1.3, a *coordination* game with $A = B$: both players drive

on a mountain road; they lose if drive on the same side of the road and win if they avoid each other, regardless of which side they take.

Figure 1.3   Coordination game.

Both (mountain,valley) and (valley,mountain) are Nash equilibria; in fact, every mixed strategy where $x_1^1 = x_2^2$, where x_i^p is the probability that player p assigns to strategy s_i^p, is a Nash equilibrium.

1.3   Some Complexity Classes

We some standard definitions of computational complexity theory; we then move on to the more recent classes **TFNP** and **PPAD**, first introduced in [11] and [16] respectively. The latter, in particular, is a key concept in the study of the problems that are the focus of this thesis.

A *Turing machine* $\mathcal{M}$ is a representation of a *program* that takes an *input*, runs a *program* manipulating the input, and either does not come to a *halting state* or it returns an *output*; the latter can be YES (in which case the Turing machine *accepts* the input), NO (the Turing machine *rejects* the input), or a string $\mathcal{M}(x)$. Formally, $\mathcal{M} = (K, \Sigma, \delta, s)$. The finite set K is the set of *states*; Σ is a finite set of *symbols* (the *alphabet* of $\mathcal{M}$) such that $\Sigma \cap K = \varnothing$, and Σ always contains the symbols $\sqcup$ (*blank*) and $\triangleright$ (*first symbol*); δ is the *transition function*

$$\delta : K \times \Sigma \longrightarrow (K \cup \{h, Yes, No\}) \times \Sigma \times \{\leftarrow, \rightarrow, -\}$$

where h is the *halting state*, YES is the *accepting state*, *No* is the *rejecting state*,

and $\{\leftarrow, \rightarrow, -\} \not\subseteq (K \cup \Sigma)$ correspond to the *cursor directions* "left," "right" and "stay."

A *language* is a set of strings of symbols $L \subset (\Sigma \setminus \{\sqcup\})^*$; a Turing machine $\mathcal{M}$ *decides* L if for every $x \in (\Sigma \setminus \{\sqcup\})^*$ $\mathcal{M}(x) = Yes$ if $x \in L$ and $\mathcal{M}(x) = No$ if $x \notin L$. We say that $\mathcal{M}$ *accepts* L if for every $x \in (\Sigma \setminus \{\sqcup\})^*$ $\mathcal{M}(x) = Yes$ if $x \in L$ and $\mathcal{M}(x)$ does not halt if $x \notin L$. Given a function $f : (\Sigma \setminus \{\sqcup\})^* \to \Sigma^*$, we say that $\mathcal{M}$ *computes* f if for every $x \in (\Sigma \setminus \{\sqcup\})^*$ $\mathcal{M}(x) = f(x)$.

Given a problem, a specific input of a problem is called an *instance*. The output of a *decision problem* P is either YES or NO. Its *complement* is the problem $\bar{P}$ that outputs "NO" for each instance of P that outputs "YES", and vice versa. A *function problem* outputs a more generic y that satisfies a binary relation $R(x, y)$, where x is the instance of the problem. *Search problems* are function problems that return either y such that $R(x, y)$, or "NO" if it's not possible to find any such y. If y is guaranteed to exist, the problem is called a *total function problem*.

An example of decision problem is: "(input) given a graph, (question) is it possible to find an Euler tour of the graph?" Its complement is "(input) given a graph, (question) is it possible that there isn't any Euler of the graph?" A search problem is: "(input) given a graph, (output) return one Euler tour of the graph, or "NO" if no such tour exists." A total function problem is: "(input) given an Euler graph, (output) return one of its Euler tours."

Let P_1 be a problem and let x be an instance of P_1 that is encoded in $|x|$ bits. P_1 *reduces to the problem P_2 in polynomial time*, denoted $P_1 \leq_P P_2$, if there exists a *polynomial-time reduction*, that is, a function $f : \{0, 1\}^* \to \{0, 1\}^*$ and a Turing machine $\mathcal{M}$ such that for all $x \in \{0, 1\}^*$

1. $x \in P_1 \iff f(x) \in P_2$;

2. $\mathcal{M}$ computes $f(x)$;

3. $\mathcal{M}$ stops after $p(|x|)$ steps, where p is a polynomial.

Intuitively, if P_1 is polynomial-time reducible to P_2, it takes polynomial time to "translate" P_1 to P_2, and then to "translate back" a solution of P_2 as a solution of P_1. This is particularly useful if P_2 is "difficult to solve"; then the problem P_1 is at least as "difficult."

A *class* is a set of languages. For any class C of decision problems, the class of all complements of the problems in C is the *complement class* co $-$ C. A problem P is *hard* for a class C if for every problem P_C in C there is a polynomial-time reduction to P; that is, if P is hard to solve at least as every problem in C. A C $-$ *hard* problem in C is *complete* for C.

The complexity class **P** contains all the *polynomially decidable problems*; that is, all problems P such that there exists a Turing machine $\mathcal{M}$ that outputs either Yes or No for all inputs $x \in \{0,1\}^*$ of P after $p(|x|)$ steps, where p is a polynomial. Intuitively, a decision problem is in **P** if the answer to its question can be found in a number of steps that is polynomial in the input of the problem. A problem P belongs to the class **NP**, *non-deterministic polynomial-time problems*, if there exists a Turing machine $\mathcal{M}$ and polynomials p_1, p_2 such that

1. for all $x \in P$ there exists a *certificate* $y \in \{0,1\}^*$ which satisfies $|y| \leq p_1(|x|)$;

2. $\mathcal{M}$ accepts the combined input xy, stopping after at most $p_2(|x| + |y|)$ steps;

3. for all $x \notin P$ there does not exist $y \in \{0,1\}^*$ such that $\mathcal{M}$ accepts the combined input xy.

Intuitively, a decision problem is in **NP** if it takes polynomial time to verify whether the "certificate solution" y is, indeed, a correct answer to the question posed by the problem. The class #**P** is the class of all problems that output the number of possible certificates for a problem in **NP**.

In [11], Megiddo and Papadimitriou introduced the classes **FNP**, *function non-deterministic polynomial*, and **TFNP**, *total function non-deterministic polynomial*. The former is defined as the class of binary relations $R(x, y)$ such that there is a polynomial-time algorithm that decides $R(x, y)$ for x, y such that $|y| \leq p(|x|)$, where p is a polynomial. The latter is the class of all such problems for which y is guaranteed to exist. Intuitively, **FNP** and **TFNP** are similar to **NP**, but they allow for problems of (respectively) function and total function form. Also in [11], Megiddo and Papadimitriou proved that, unless $\mathbf{NP} = \mathbf{co} - \mathbf{NP}$, **TFNP** is a *semantic* class, that is, a class without complete problems. To circumvent this limitation of **TFNP**, Papadimitriou [16] focused on the problems for which the existence of a solution is proved by a "parity argument", introducing the classes **PPA** (*Proof by Parity Argument*) and **PPAD** (*Proof by Parity Argument, Directed version*).

Intuitively, **PPA** is the class of problems for which the existence of a solution can be proved using the argument "in any undirected graph with one odd-degree node there must be another odd-degree node." It is interesting to note that no **PPA** problems have been proven **PPA**-complete. The existence of the problems in **PPAD**, on the other hand, can be proven using the argument "in any directed graph with one unbalanced node (that is, with outdegree different from its indegree) there must be another unbalanced node." The latter can simplified without loss of generality or computational power to the case of indegree and outdegree at most one, that is, "in any directed graph in which all vertices have indegree and outdegree at most one, if there is a *source* (a node with indegree zero), then there must be a *sink* (a node with outdegree zero)." Formally, we can define **PPAD** as the class of problems reducible to the problem END OF THE LINE.

input : Two circuits S and P with n input bits and n output bits such that $P(0^n) = 0^n \neq S(0^n)$.

output: An input $x \in \{0,1\}^n$ such that $P(S(x) \neq x)$ or $S(P(x)) \neq x$ and $S(P(x)) \neq 0^n$

This is the definition given in Daskalakis, Goldberg and Papadimitriou [5]; Papadimitriou [16] defines the class in terms of Turing machines. A circuit is formally defined as a directed acyclic graph with n vertices with indegree 0 called *input nodes*, m vertices with outdegree 0 called *output nodes*, and *internal nodes* with indegree 1 or 2; when each input node receives an input in $\{0,1\}$, the internal nodes with indegree 2 compute the Boolean functions *and* or *or*, the internal nodes with indegree 1 compute the Boolean function *not* and each output node returns a value in $\{0,1\}$ accordingly. The problems in **PPAD** can be seen as a circuit S ("successor"), and a circuit P ("predecessor") that are used to build a directed graph with an edge (x, y) if and only if $S(x) = y$ and $P(y) = x$, with a *standard source* 0^n; the output is either a sink or a non-standard source. Problems in **PPA** can be defined analogously, but the resulting graph is not directed and instead of sources and sinks there are generic endpoints; that is, the roles of P and S are interchangeable. We have that PPAD $\subset$ PPA, and there are **PPA** problems that are not **PPAD**, as we will see in section 2.1. Figure 1.4 presents an example of graph of a **PPAD** problem; a graph for a **PPA** problem is analogous, but undirected.

By theorem 1, the problem n-NASH, defined as follows, is a total function problem.

n-NASH

input : A n-player game.

output: A Nash equilibrium of the game.

Figure 1.4   In green, the circuit S; in red, the circuit P. The standard source is the black node; the green nodes are the other sources; the red nodes are the sinks. The graph can include paths, cycles and isolated points.

Megiddo and Papadimitriou ([11]) proved that it is in **TFNP**. Daskalakis, Goldberg and Papadimitriou [5] and Chen and Deng [4] have proven its **PPAD**-completeness, the former for $n \geq 3$ and the latter for $n \geq 2$. A small amendment of the proof in [5] can be found in Casetti [2].

Theorem 2. (Daskalakis, Goldberg and Papadimitriou [5]; Chen and Deng [4]) *For $n \geq 2$, the problem n-NASH is* **PPAD***-complete.*

We will see more problems in **PPA** and **PPAD** in chapter 2; in fact, our main result can be seen as a negative result on the **PPAD** complexity of a case of 2-Nash.

1.4 Bimatrix Games and Labels

In the rest of this thesis we will focus bimatrix games. We will assume that the payoff matrices (A, B) are non-negative, and that neither A nor $B^\top$ has a zero column, if necessary applying an affine transformation, which does not affect the equilibria of the game.

The Nash equilibria of bimatrix games can be analysed from a combinatorial point of view using *labels*. This method is due to Shapley [19], in a study building on ideas introduced in a paper by Lemke and Howson [9]. Let $n, m \in \mathbb{N}$ with $m \leq n$, and consider a set X with $|X| = n$. A *labeling* of X is a function $l : X \to [m]$. An m-uple $x = (x_1, \ldots, x_m) \in X^m$ is *completely labeled* if $\{i \in [m] \mid l(x_j) = i \text{ for some } j \in [m]\} = [m]$, that is, if each label $j \in [m]$ appears once and only once in $(l(x_1), \ldots, l(x_m))$.

Let (A, B) be bimatrix game, and let X and Y be the mixed strategy simplices of respectively player 1 and 2; that is

$$X = \{x \in \mathbb{R}^m \mid x \geq \mathbf{0}, \ \mathbf{1}^\top x = 1\}; \quad Y = \{y \in \mathbb{R}^n \mid y \geq \mathbf{0}, \ \mathbf{1}^\top y = 1\}. \quad (1.1)$$

A *labeling* of the game is then given as follows:

1. the m pure strategies of player 1 are identified by $1, \ldots, m$;

2. the n pure strategies of player 2 are identified by $m + 1, \ldots, m + n$;

3. each mixed strategy $x \in X$ of player 1 has

   - label i for each $i \in [m]$ such that $x_i = 0$, that is if in x player 1 does not play her ith pure strategy,

   - label $m + j$ for each $j \in [n]$ such that the jth pure strategy of player 2 is a best response to x;

4. each mixed strategy $y \in Y$ of player 2 has

   - label $m + j$ for each $j \in [n]$ such that $y_j = 0$, that is if in y player 2 does not play his jth pure strategy,

- label i for each $i \in [m]$ such that the ith pure strategy of player 1 is a best response to y.

The labeling of mixed strategy profiles can be used to characterize the Nash equilibria of the game.

Theorem 3. (Shapley [19]) *Let $(x, y) \in X \times Y$; then (x, y) is a Nash equilibrium of the bimatrix game (A, B) if and only if (x, y) is completely labeled.*

Proof. The mixed strategy $x \in X$ has label $m + j$ for some $j \in [n]$ if and only if the jth pure strategy of player 2 is a best response to x; this, in turn, is a necessary and sufficient condition for player 2 to play his jth strategy at an equilibrium against x. Therefore, at an equilibrium (x, y) all labels $m + 1, \ldots, m + n$ will appear either as labels of x or of y. The analogous holds for the labels $i \in [n]$. $\qquad\qquad \square$

An useful graphical representation of labels on the simplices X and Y. is done by labeling the outside of each simplex according to the player's own pure strategies that are *not* played, and by subdividing its interior in closed polyhedral sets corresponding to the other player's best response, called *best response regions*. We give an example of this construction.

Example 1.2. Consider the 3×3 game (A, B) with

$$A = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}, \qquad B = \begin{pmatrix} 0 & 2 & 4 \\ 3 & 2 & 0 \\ 0 & 2 & 0 \end{pmatrix}. \qquad (1.2)$$

The pure strategies of player 1 are labeled as $1, 2, 3$; the pure strategies of player 2 are labeled as $4, 5, 6$.

In the following figures the labels of the strategies will be represented as circled numbers. Figure 1.5 shows X and Y: the exterior is labeled with the player's own pure strategies where these are not played, so the facet labeled with a pure strategy is opposite to the vertex where only that pure strategy

is played; the interior is covered by the best response regions, with labels corresponding to the other player's pure strategies that are in the best response to the player's own strategy in the region. For example, the best-response region in Y with label 1 is the set of those (y_1, y_2, y_3) so that $y_1 \geq y_2$ and $y_1 \geq y_3$. There is only one pair (x, y) that is completely labeled, namely $x = (\frac{1}{3}, \frac{2}{3}, 0)$ with labels $3, 4, 5$ and $y = (\frac{1}{2}, \frac{1}{2}, 0)$ with labels $1, 2, 6$, so this is the only Nash equilibrium of the game.

Figure 1.5 Labeled mixed strategy sets X and Y for the game (1.2).

The point of view of best response regions can be translated to an equivalent construction of *best response polytopes*. We begin by noticing that the best-response regions can be obtained as projections on X and Y of the *best-response facets* of the polyhedra

$$\bar{P} = \{(x, v) \in X \times \mathbb{R} \mid B^\top x \leq \mathbf{1}v\}; \quad \bar{Q} = \{(y, u) \in Y \times \mathbb{R} \mid Ay \leq \mathbf{1}u\}. \quad (1.3)$$

These facets in $\bar{P}$ are defined as the points $(x, v) \in X \times \mathbb{R}$ such that $(B^\top x)_j = v$, which in turn represent the strategies $x \in X$ of player 1 that give exactly payoff v to player 2 when he plays strategy j; the projection of the facet defined by $(B^\top x)_j = v$ to X has then label j. Analogously, the facet of $\bar{Q}$ given by the points $(y, u) \in Y \times \mathbb{R}$ such that $(Ay)_i = u$ will project to the best-response region of Y with label i.

Example 1.3. In the example 1.2, the inequalities $B^\top x \leq \mathbf{1}v$ state $3x_2 \leq v$, $2x_1 + 2x_2 + 2x_3 \leq v$, $4x_1 \leq v$. Figure 1.6 shows the best-response facets of

$\bar{P}$ and their projection to X by ignoring the payoff variable v, which gives the subdivision of X into best-response regions Figure 1.5.

Figure 1.6 Best response facets of the polyhedron $\bar{P}$ in (1.3).

Given the assumptions on non-negativity of A and $B^\top$, we can give a change coordinates to x_i/v and y_j/u and replace $\bar{P}$ and $\bar{Q}$ with the *best-response polytopes*

$$P = \{x \in \mathbb{R}^m \mid x \geq \mathbf{0}, \ B^\top x \leq \mathbf{1}\}; \quad Q = \{y \in \mathbb{R}^n \mid y \geq \mathbf{0}, \ Ay \leq \mathbf{1}\}. \quad (1.4)$$

The polytope P is the intersection of $n + m$ half spaces, one for each inequality corresponding to either player 1 avoiding her i-th pure strategy, where $i \in [m]$, or to a best response of player 2 that gives non-zero probability to his j-th strategy, where $j \in [n]$. Formally, a point $x \in P$ has label k if and only if either $x_k = 0$ for $k \in [m]$ or $(B^\top x)_k = 0$ for $k \in [n]$. Analogously, a point in Q has label k if and only if either $y_k = 0$ for $k \in [n]$ or $(Ay)_k$ for $k \in [m]$. Then a point $(x, y) \in P \times Q$ is completely labeled if and only if it satisfies the *complementarity condition*

$$x_i = 0 \text{ or } (Ay)_i = 1 \text{ for all } i \in [m];$$
$$y_j = 0 \text{ or } (B^\top x)_j \text{ for all } j \in [n]. \quad (1.5)$$

Then either he point corresponding to (x, y) in $\bar{P} \times \bar{Q}$ is a Nash equilibrium, or $(x, y) = (\mathbf{0}, \mathbf{0})$; we will refer to the latter case as *artificial equilibrium*.

Example 1.4. Keeping on with example 1.2, the polyhedron $\bar{P}$ of Figure 1.6 becomes the polytope in Figure 1.7 1.7.

Figure 1.7 Best response polytope P in (1.4) for the game in (1.2).

Note that the vertex $(\mathbf{0}, \mathbf{0})$ has labels $1, 2, 3$ in P and, analogously, it has labels $4, 5, 6$ in Q, so it is completely labeled, and corresponds to the "artificial" equilibrium.

We will now see some special cases of games and a series of polynomial-time reductions that connect them, so that a method that finds an equilibrium of these games can be used to find an equilibrium of any arbitrary bimatrix game.

First of all, we focus on symmetric games. Any bimatrix game can be "symmetrized"; the result is due to Gale, Kuhn and Tucker [8] for zero-sum games; its extension to non-zero-sum games is a folklore result.

Proposition 1.2. *Let* (A, B) *be a bimatrix game and* (x, y) *be one of its Nash equilibria. Then* (z, z), *where* $z = (x, y)$, *is a Nash equilibrium of the symmetric game* $(C, C^\top)$, *where*

$$C = \begin{pmatrix} 0 & A \\ B^\top & 0 \end{pmatrix}. \tag{1.6}$$

It takes polynomial time in the size of a matrix to calculate its transpose; the rest of the reduction from (A, B) to $(C, C^\top)$ is linear. Therefore, we have a reduction from 2-NASH to SYMMETRIC NASH.

McLennan and Tourky [10] have proven a result in the opposite direction of proposition 1.2: any symmetric game can be translated into a bimatrix game of the form (I, B), called *imitiation game*.

Theorem 4. (McLennan and Tourky [10]) *The pair (x, x) is a symmetric Nash equilibrium of the symmetric bimatrix game $(C, C^\top)$ if and only if there is some y such that (x, y) is a Nash equilibrium of the imitation game (I, B) with $B = C^\top$.*

Again, the reduction from $(C, C^\top)$ takes polynomial time (even better: linear). So we have a reduction from SYMMETRIC NASH to IMITATION NASH.

We see now an example illustrating theorem 4. Note how the mixed strategy x of player 1 in any Nash equilibrium of (I, B) corresponds exactly to the symmetric equilibrium (x, x) in the symmetric game defined by the payoff matrix of player 2. Note also how theorem 4 applies to the symmetric equilibria of the symmetric game, but not to all its equilibria; there could be non-symmetric equilibria of $(C, C^\top)$ that are not found through the imitation game.

*Example* 1.5. As an example, consider the symmetric game $(C, C^\top)$ with

$$C = \begin{pmatrix} 0 & 3 & 0 \\ 2 & 2 & 2 \\ 4 & 0 & 0 \end{pmatrix}, \qquad C^\top = \begin{pmatrix} 0 & 2 & 4 \\ 3 & 2 & 0 \\ 0 & 2 & 0 \end{pmatrix}, \tag{1.7}$$

so that $C^\top = B$ in (1.2), and the corresponding imitation game $(I, C^\top)$ is (A, B) in (1.2).

Figure 1.8 shows the labeled mixed-strategy simplices X and Y for the game 1.7; since the game is symmetric, only the labels are different. In addition to the symmetric equilibrium (x, x) where $x = (\frac{1}{3}, \frac{2}{3}, 0)$, the game has two non-symmetric equilibria in (a, b) and (b, a) with $a = (\frac{1}{2}, \frac{1}{2}, 0)$ and $b = (0, \frac{2}{3}, \frac{1}{3})$.

Note that the corresponding imitation game (A, B) has only one equilibrium in (x, y), where (x, x) is the symmetric equilibrium of $(C, C^\top)$.

Figure 1.8 Best response regions for the symmetric game $(C, C^\top)$ in (1.7).

The characterization of Nash equilibria as completely labeled pairs (x, y) holds for arbitrary bimatrix games. From now on, we will also impose a further condition: that the games are *nondegenerate* This means that all points in P have at most m labels, and all points in Q have at most n labels. Therefore, in an equilibrium (x, y) each label appears exactly once. This also means that the number of pure best response strategies against a mixed strategy is never larger than the size of the support of that mixed strategy. Geometrically, no point of P lies on more than m facets and no point of Q lies on more than n facets, so the best response polytopes P and Q are simple. Furthermore, a point of P has exactly m labels if and only if it is a vertex, and a point of Q has exactly n labels if and only if it is a vertex; therefore, all completely labeled points (x, y) are vertices, and Nash equilibria are isolated points.

Example 1.6. An example of degenerate game is given by $(C, C^\top)$ where

$$C = \begin{pmatrix} 0 & 4 & 0 \\ 2 & 2 & 2 \\ 4 & 0 & 0 \end{pmatrix}, \qquad C^\top = \begin{pmatrix} 0 & 2 & 4 \\ 4 & 2 & 0 \\ 0 & 2 & 0 \end{pmatrix}. \tag{1.8}$$

As it is shown in Figure 1.9, the game $(C, C^\top)$ is degenerate because the mixed strategy $x = (\frac{1}{2}, \frac{1}{2}, 0)$, that also defines the unique symmetric equilibrium (x, x) of the game, has three pure best responses.

Note that the corresponding equilibria (x, y) of the imitation game $(I, C^\top)$

Figure 1.9 Best-response regions of the degenerate symmetric game (1.8) with a unique symmetric equilibrium.

are not unique, because due to degeneracy any convex combination of $(\frac{1}{2}, \frac{1}{2}, 0)$ and $(\frac{1}{3}, \frac{1}{3}, \frac{1}{3})$ can be chosen for y, as shown in Figure 1.10.

Figure 1.10 Labeled mixed-strategy sets for the imitation game $(I, C^{\top})$ for the degenerate symmetric game (1.8) where the equilibria (x, y) are not unique.

This shows how the reduction between symmetric equilibria (x, x) of a symmetric game and Nash equilibria (x, y) of the corresponding imitation game as in theorem 4 does not preserve uniqueness if the game is degenerate.

A generalization of imitation games is the class of *unit vector games*, that is games of the form (U, B) where the columns of the matrix U are unit vectors. By the results above, finding a Nash equilibrium of a bimatrix game is at least as hard as finding a Nash equilibrium of a unit vector game; that is, 2-Nash

26

reduces to Unit Vector Nash. In unit vector games, the problem of finding a completely labeled vertex of the polytope $P \times Q$ can be translated into te problem o finding a completely labeled vertex of one simple polytope that encodes all the game.

The result was due by Balthasar [1] in a different form and given here as in Savani and von Stengel [18].

Theorem 5. (Savani and von Stengel [18]) *Let* $l : [n] \to [m]$, *and let* (U, B) *be the unit vector game where* $U = (e_{l(1)} \ \cdots \ e_{l(n)})$. *Let* $N_i = \{j \in [n] \mid l(j) = i\}$ *for* $i \in [m]$, *and consider the polytopes* P^l *and* Q^l

$$P^l = \{x \in \mathbb{R}^m \mid x \geq \mathbf{0}, \ B^\top x \leq \mathbf{1}\}; \tag{1.9}$$

$$Q^l = \{y \in \mathbb{R}^n \mid y \geq \mathbf{0}, \ \sum_{\substack{j \in N_i \\ i \in [m]}} y_j \leq 1\}. \tag{1.10}$$

Let l_f *be the labeling of the facets of* P^l *defined as follows:*

$$x_i \geq 0 \ \text{has label } i \ \text{for } i \in [m];$$

$$(B^\top x)_j \leq 1 \ \text{has label } l(j) \ \text{for } j \in [n]. \tag{1.11}$$

Then $x \in P^l$ *is a completely labeled vertex of* $P^l \setminus \{\mathbf{0}\}$ *if and only if there is some* $y \in Q^l$ *such that, after scaling, the pair* (x, y) *is a Nash equilibrium of* (U, B)

Proof. Let P, Q be the polytopes associated to the game (U, B) as before.

Let $(x, y) \in P \times Q \setminus \{\mathbf{0}, \mathbf{0}\}$ be a Nash equilibrium of (U, B), therefore completely labeled in $[m + n]$. If $x_i = 0$, then x has label $i \in m$. If $x_i > 0$, then y has label i, so $(Uy)_i = 1$; then for some $j \in [n]$ we have $y_j > 0$ and $U_j = e_i$; that is, we have $y_j > 0$ and $l(j) = i$ for some $j \in [n]$. Since $y_j > 0$, $x \in P$ has label $m + j$; then, $(B^\top x)_j = 1$;s therefore $x \in P^l$ has label $l(j) = i$. Hence, x is a completely labeled vertex of P^l.

Conversely, let $x \in P^l \setminus \{\mathbf{0}\}$ be completely labeled. If $x_i > 0$, then there is $j \in [m]$ such that $(B^\top x) = j$ and $l(j) = i$, that is, $j \in N_i$. For all i such

that $x_i > 0$, define y as follows: $y_j = 1$; $y_h = 0$ for all $h \in N_i \setminus \{j\}$. Then $(x, y) \in P \times Q$ is completely labeled. $\qquad\qquad\qquad\qquad\qquad\qquad\square$

Example 1.7. The game in example 1.2 is a unit vector game, with $l(i) = i$. In figure 1.7 we have the labels $4, 5, 6$ on the best response facets of polytope P; the facets with labels $1, 2, 3$, that is, the facets where $x_1 = 0$, $x_2 = 0$, $x_3 = 0$, respectively, are at back right, back left, and bottom of the polytope. In the polytope P^l the labels $4, 5, 6$ are replaced by $1, 2, 3$ because the corresponding columns of A are the unit vectors e_1, e_2, e_3. Figure 1.11 shows the polytope P^l; the only hidden facet, with label 1, is at the bottom. Apart from the origin $\mathbf{0}$, the only completely labeled point of P^l is x, as shown in figure 1.11.

Figure 1.11 The polytope P^l for the unit vector game (1.2). The hidden facet at the back has label 1.

We will now move on the dual version given of theorem 5, as ginve in Balthasar [1]. We traslate the polytope P^l as in 1.9 to $P = \{x - \mathbf{1} \mid x \in P^l\}$, and multiply all payoffs in B by a constant, if necessary, so that $\mathbf{0}$ is in the interior of P. We have that

$$P = \{x + \mathbf{1} \geq \mathbf{0}, \ (x + \mathbf{1})^\top B \leq \mathbf{1}\} =$$
$$= \{x \in \mathbb{R}^m \mid \ -x_i \leq 1 \text{ for } i \in [m], \ x^\top (b_j/(1 - \mathbf{1}^\top b_j)) \leq 1 \text{ for } j \in [n]\}.$$

The polar of P is then

$$P^\Delta = \text{conv}(\{e_i \mid i \in [m]\} \ \cup \ \{c_j \mid j \in [n]\}) \qquad\qquad (1.12)$$

with $c_j = b_j/(1 - \mathbf{1}^\top b_j)$. Since P is a simple polytope with $\mathbf{0}$ in its interior, P^Δ is simplicial, $P^{\Delta\Delta} = P$, and the facets of P^Δ correspond to the vertices of P and vice versa. We can then label the vertices of P^Δ as the corresponding facets in P^l; the completely labeled facets of P^Δ will then correspond to the completely labeled vertices of P^l. In particular, the facet corresponding to $\mathbf{0}$ will be

$$F_0 = \{x \in P^\Delta \mid -\mathbf{1}^\top x = 1\} = \operatorname{conv}\{e_i \mid i \in [m]\}. \tag{1.13}$$

Theorem 5 then translates as in the original version by Balthasar [1].

Theorem 6. (Balthasar [1]) *Let Q be a labeled m-dimensional simplicial polytope with $\mathbf{0}$ in its interior and vertices $e_1, \ldots, e_m, c_1, \ldots, c_n$ such that $F_0 = \operatorname{conv}\{e_i \mid i \in [m]\}$ is a facet of Q. Let (U, B) be a unit vector game, with $U = (e_{l(1)} \; \cdots \; e_{l(n)})$ for a labeling $l : [n] \to [m]$ and $B = (b_1 \; \cdots \; b_n)$ with $b_j = c_j/(1 + \mathbf{1}^\top c_j)$ for $j \in [n]$. Let l_v be a labeling of the vertices of Q as follows:*

$$l_v(-e_i) = i \text{ for } i \in [m];$$
$$l_v(c_j) = l(j) \text{ for } j \in [n]. \tag{1.14}$$

Then a facet $F \neq F_0$ of Q with normal vector v is completely labeled if and only if (x, y) is a Nash equilibrium of (U, B), where $x = (v + \mathbf{1})/(\mathbf{1}^\top(v + \mathbf{1}))$, so $x_i = 0$ if and only if $e_i \in F$ for $i \in [m]$, and the mixed strategy y is the uniform distribution on the set of the pure best replies to x, which in turn correspond to $j \in [n]$ such that c_j is a vertex of F.

As in theorem 5 we have a correspondence between completely labeled vertices of P^l and equilibria of the unit vector game (U, B) with an "artificial" equilibrium corresponding to the vertex $\mathbf{0}$, in theorem 6 we have a correspondence between the completely labeled facets of the polytope Q and the equilibria of the unit vector game (U, B) with the "artificial" equilibrium corresponding to the facet F_0.

Given a bimatrix nondegenerate game (A, B), it takes polynomial time to write and solve the linear equations defining its best response polyhedra $\bar{P}, \bar{Q}$ and its best response polytopes P, Q. It also take polynomial time to label $\bar{P}, \bar{Q}$ and P, Q. Analogously, given a unit vector game (U, B), it takes polynomial time to construct and label the polytope P^l. Therefore, by theorem 5, we have a polynomial time reduction from the problem 2-Nash for nondegenerate games to the problem Another Completely Labeled Vertex, defined as follows.

Another Completely Labeled Vertex

input : A simple m-dimensional polytope P with $m + n$ facets; a labeling
$\quad\quad\quad l_f : [m + n] \to [n]$; a facet F_0 of P, completely labeled by l_f.

output: A facet $F \neq F_0$ of S, completely labeled by l.

Proposition 1.3. *The problem* Nondegenerate 2-Nash *reduces in polynomial time to* Another Completely Labeled Vertex.

Theorem 6 gives the dual of proposition 1.3; since the construction of the polar polytope from the original one is also polynomial, the reduction is to the problem Another Completely Labeled Facet.

Another Completely Labeled Facet

input : A simplicial m-dimensional polytope Q with $m + n$ vertices; a
$\quad\quad\quad$ labeling $l_v : [m + n] \to [n]$; a facet F_0 of Q, completely labeled by
$\quad\quad\quad l_v$.

output: A facet $F \neq F_0$ of S completely labeled by l_v.

Proposition 1.4. *The problem* Nondegenerate 2-Nash *reduces in polynomial time to* Another Completely Labeled Facet.

1.5 Cyclic Polytopes and Gale Strings

In theorems 5 and 6 we have built a correspondence between labeled polytopes and unit vector games, where Nash equilibria correspond to completely labeled vertices or facets. We now focus on a particular kind of simplicial polytopes, called *cyclic polytopes*, that can be represented as a combinatorial structure, the *Gale strings*. We will first give the definition of cyclic polytope, then of Gale string, then the theorem by Gale [7] about their correspondence.

The *moment curve* in dimension d is defined as

$$\mu_d : \mathbb{R} \longrightarrow \mathbb{R}^d, \quad \mu_d : t \longmapsto (t, t^2, \ldots, t^d)^\top. \tag{1.15}$$

The *cyclic polytope* in dimension d with n vertices, where $n > d$ is

$$C_d(n) = \text{conv}\{\mu_d(t_i) \mid t_1 < \ldots < t_n \text{ affinely independent }\}. \tag{1.16}$$

Example 1.8. The cyclic polytope in dimension 3 with 6 facets can be seen in figure 1.12.

Figure 1.12 The cyclic polytope $C_3(6)$

Given $k \in \mathbb{N}$ and a set S, we can represent the function $f : [k] \to S$ as the string $s = s(1)s(2)\cdots s(k)$; we have a *bitstring* if $S = \{0, \mathbf{1}\}$. A maximal substring of consecutive $\mathbf{1}$'s in a bitstring is called a *run*; an *interior* run is bounded on both sides by 0's. We will use the notation $\mathbf{1}^k$ for a run of length

31

k, and 0^k for a string of 0's of length k A *Gale string of length n and dimension d*, where $n > d$, is a bitstring $s \in G(d, n)$ satisfying the following conditions:

1. exactly d bits in s are **1** and

2. (*Gale evenness condition*)

$$01^k0 \text{ is a substring of } s \implies k \text{ is even.} \qquad (1.17)$$

In general, the Gale evenness conditions allows for Gale strings that start or end with an odd-length run; but if d is even then s can start with an odd run if and only if it ends with an odd run. We can then consider the Gale strings in $G(d, n)$ with even d as a the "loops" obtained by "glueing together" the extremes of the strings, so that all runs on the loops are even. Formally: we can see the indices of a Gale string $s \in G(d, n)$ with d even as equivalence classes modulo n, identifying $s(i + n) = s(i)$. This also shows that the set of Gale strings of even dimension is invariant under a cyclic shift of the strings.

Example 1.9. As an example of d even, we have

$$G(4, 6) = \{\mathbf{111100}, \mathbf{111001}, \mathbf{110011}, \mathbf{100111}, \mathbf{001111},$$

$$\mathbf{011110}, \mathbf{110110}, \mathbf{101101}, \mathbf{011011}\}$$

The strings **111100**, **111001**, **110011**, **100111**, **001111** and **011110** are equivalent under a cyclic shift (if considering the strings as "loops", the **1**'s are all consecutive), as are the strings **110110**, **101101** and **011011** (if considering the strings as "loops", the even runs of **1**'s are two couples separated by a single 0).

As an example for d odd, we have

$$G(3, 5) = \{\mathbf{11100}, \mathbf{10110}, \mathbf{10011}, \mathbf{11001}, \mathbf{01101}, \mathbf{00111}\}$$

Note how **01011** is a cyclic shift of **10110**, but it is not a Gale string.

The relation between cyclic polytopes and Gale strings is given by the following theorem by Gale [7].

Theorem 7. (Gale [7]) *For any positive integers d, n with $n > d$*

$$F \text{ is a facet of } C_d(n)$$

$$\Longleftrightarrow$$

$$F = \text{conv}\{\mu(t_j) \mid s(j) = 1 \text{ for } s \in G(d, n)\}. \tag{1.18}$$

Proof. We have that $C_d(n) = \text{conv}\{\mu_d(\bar{t}_j) \mid \bar{t}_1 < \ldots < \bar{t}_n \text{ affinely independent}\}$. We choose $t_1 < \ldots < t_d$ in the $\bar{t}_j$'s as above; then the points $\mu_d(t_i)$ with $i \in [d]$ define an hyperplane H that crosses the moment curve $\mu_d(t)$ at all and only the t_i's with $i \in [d]$. The hyperplane H is never tangent to the moment curve, and every crossing gives a "change of sign"; that is, if there is one and only one $t_i \in (t, t')$ we have that $(\mu_d(t))(\mu_d(t')) < 0$.

A facet F of the cyclic polytope $C_d(n)$ corresponds to a choice of t_i's with $i \in [d]$ such that for all the $\bar{t}_k$'s with $k \in [n]$ and $\bar{t}_k \notin \{t_i \mid i \in [d]\}$ all the $\mu_d(\bar{t}_k)$'s have the same sign. This can happen only if the moment curve has an even number of changes of sign, therefore of t_i's, between two $\bar{t}_k$'s. The t_i's will correspond to the **1**'s of the Gale string $s \in G(d, n)$ and the $\bar{t}_k$'s to the 0's, and the condition for being a facet becomes the Gale evenness condition. $\square$

Note also that the fact that the moment curve has exactly d zeroes implies that the each facet of $C_d(n)$ is a d-simplex, so $C_d(n)$ is simplicial; furthermore, the choice of the $\bar{t}_j$'s in the proof is irrelevant.

*Example* 1.10. Consider the facet F of the cyclic polytope $C_3(6)$ underlined in blue in figure 1.16.

If we label the vertices on the moment curve as $i \in [n]$, and we set $i = 1$ if i is a vertex of F and $i = 0$ otherwise, we see that the corresponding Gale string $s \in G(3, 6)$ is $s = \mathbf{1}00\mathbf{11}0$.

On the hyperplane, the correspondence is as in figure 1.14.

Example 1.11. Figure 1.15 shows the cyclic polytope $C_4(6)$, with vertices labeled according to the index $i \in [6]$ of the Gale string. The thin lines represent

Figure 1.13   A facet of the cyclic polytope $C_3(6)$.

Figure 1.14   A facet of the cyclic polytope $C_3(6)$ and the corresponding Gale string
$s \in G(3,6)$.

edges inside the facet 1234, drawn in bold lines, corresponding to the Gale
string $s = \mathbf{1111}00$.

On the hyperplane, the string $s = \mathbf{1111}00$ can be seen as in figure **??**.

Example 1.12. As a counterexample, consider figure 1.17. The corresponding
bitstring would be $s = \mathbf{111}0\mathbf{1}0$, not a Gale string. The point in pink and the
point in green are on different sides of the hyperplane H, so they do not belong
to the same facet.

Analogously for the string $s = \mathbf{1}0\mathbf{1}0\mathbf{1}$, corresponding to figure 1.18

We want now to exploit theorem 11 to study bimatrix games; more specif-
ically, unit vector games. We have seen in proposition 1.4 that 2-NASH can
be reduced to ANOTHER COMPLETELY LABELED FACET. If the polytope Q

Figure 1.15 The cyclic polytope $C_4(6)$.

is cyclic and we define a labeling for Gale strings such that a completely labeled Gale string corresponds to a completely labeled facet of the polytope Q in theorem 6, we can study unit vector games with dual cyclic best response polytope as Gale strings. Note that it takes polynomial time to translate a cyclic polytope $C_d(n)$ into the corresponding $G(d, n)$: the bitstrings of length n with d positive bits are found in $O(n^2)$ time, and checking that a bitstring satisfies the Gale evenness condition takes $O(d)$ time, so the operation takes

Figure 1.16 The Gale string $s = \mathbf{111100}$.

Figure 1.17 The points on the moment curve have different signs, so they don't correspond to a facet; the corresponding bitstring does not satisfy the Gale evenness condition.

Figure 1.18 The points on the moment curve don't correspond to a facet, and the corresponding bitstring does not satisfy the Gale evenness condition.

$O(n^3 d)$ time.

We say that $s \in G(d, n)$ is a *completely labeled Gale string* if for some labeling function $l_s : [n] \to [d]$ the set $\{i \in [n] \mid s(i) = \mathbf{1}\}$ is completely labeled by l_s. Since $s \in G(d, n)$ has exacty d bits equal to $\mathbf{1}$, this means that for each $j \in [d]$ there is exactly one $i \in [n]$ such that $s(i) = \mathbf{1}$ and $l_s(i) = j$.

Note that it is not always possible to find a completely labeled Gale string.

Example 1.13. For $l = 121314$, there are no completely labeled Gale strings.

The labels $l(i) = 2, 3, 4$ appear only once in l, as $l(2), l(4), l(6)$ respectively; therefore we must have $s(2) = s(4) = s(6) = 1$. For every other $i \in [n]$ we

have $l(i) = 1$, so we have $l(i) = 1$ for exactly one $i = 1, 3, 5$. The candidate strings are then **110101**, **011101**, **010111**; but none of these satisfies the Gale evenness condition.

Let (U, B), where $U = (e_{l(1),...,e_{l(d)}})$ for some labeling $l : [n] \to [d]$, be a unit vector game for which the dual of the best response polytope is a cyclic polytope $Q = \text{conv}\{e_1, \ldots, e_d, c_1, \ldots, c_n\}$. Theorem 6 gives a labeling l_v of the $d + n$ vertices of Q as in 1.14:

$$l_v(-e_i) = i \text{ for } i \in [m];$$
$$l_v(c_j) = l(j) \text{ for } j \in [n].$$

We define the labeling $l_s : [d + n] \to [d]$ as follows:

$$l_s(i) = i \text{ for } i \in [d];$$
$$l_s(d + j) = l(j) \text{ for } j \in [n]. \tag{1.19}$$

Then the Gale strings $s \in G(d, d + n)$ that are completely labeled for l_s correspond exactly to the completely labeled facets of Q, with the facet F_0 corresponding to the "trivial" completely labeled string $\mathbf{1}^d 0$.

From this point forward, we will assume that d is even. We will also assume that the labeling $l : [n] \to [d]$ is such that $l(i) \neq l(i + 1)$; this can be done without loss of generality, given the following consideration. Suppose that $l(i) = l(i + 1)$ for some index i, and let s be a completely labeled Gale string for l. Then only one of $s(i)$ and $s(i + 1)$ can be equal to $\mathbf{1}$ (note that it's possible that both are 0s). So $s(i)s(i + 1)$ will never be a run of even length that "interferes" with the Gale Evenness Condition, so we can "simplify" by identifying the indices i and $i + 1$.

Example 1.14. Given the string of labels $l = 123432$, there are four associated completely labeled Gale strings in $G(4, 6)$: $s_A = \mathbf{111100}$, $s_B = \mathbf{110110}$, $s_C = \mathbf{100111}$ and $s_D = \mathbf{101101}$.

facet	1 2 3 4 3 2
A	1 1 1 1 · ·
B	1 1 · 1 1 ·
C	1 · · 1 1 1
D	1 · 1 1 · 1

In the cyclic polytope $C_4(6)$, these correspond to the facets as in figure 1.19.

Figure 1.19 The cyclic polytope $C_4(6)$; the vertices are labeled following the labeling $l_s = 123432$. The completely labeled facets correspond to the completely labeled Gale strings.

We can now define the problem ANOTHER GALE as follows:

ANOTHER GALE

---

input : A labeling $l : [n] \to [d]$, where d is even and $d < n$. A Gale string $s \in G(d, n)$, completely labeled by l.

output: A Gale string $s' \in G(d, n)$, completely labeled by l, such that $s' \neq s$.

Note that building the labeling l_s from the labeling l_v takes polynomial time: for the labels $i \in [m]$ it is immediate, for the following $n - m$ labels we have to check the labeling $l : [n] \to [m]$, that is, the $n \times m$ unit matrix of the imitation game. Therefore, we have a reduction from 2-NASH to ANOTHER GALE.

Proposition 1.5. *The problem 2-*NASH *is polynomial-time reducible to the problem* ANOTHER GALE.

Chapter 2

Algorithmic and Complexity Results

2.1 The Lemke-Howson Algorithm

In the previous chapter we have defined some problems of the form "find another completely labeled..." for vertices, facets and Gale strings. In this section we will focus on different versions of a standard algorithm, first introduced by Lemke and Howson in [9], that yields a solution to these problems using two main concepts: the *pivoting* routine and *almost complete labeled* vertices, facets and Gale strings.

Let P be a simple polytope in dimension d with n facets. We *pivot on the vertices* of P by moving from a vertex x to another vertex y connected to x by an edge. Note that, since P is simple, there are exactly d possible choices for y. Analogously, we *pivot on the facets* of a simplicial polytope Q in dimension d by moving from a facet F to a facet G that share all vertices but one; and since Q is simplicial, there are d possible choices for G.

Suppose now that there is a labeling $l_f : [n] \to [d]$ of the facets of the simple polytope P. If we pivot from a vertex x to a vertex x' we "leave behind" a facet F with label k; so, if x has labels $(l_1, \ldots, k, \ldots, l_d)$, then x' has labels

$(l_1, \ldots, h, \ldots, l_d)$, where h is the label of the facet F' that does not have x as its vertex. We call this *dropping label k and picking up label h*, or *pivoting on label k*. Analogously, if there is a labeling $l_v : [n] \to [d]$ of the vertices of the simplicial poytope Q and we pivot from a facet F with labels $(l_1, \ldots, k, \ldots, l_d)$ to a facet F' with labels $(l_1, \ldots, h, \ldots, l_d)$, we say that we *drop label k and pick up label h*, or that we *pivot on label k*.

Let $m, n \in \mathbb{N}$ with $m \leq n$; consider a set X with $|X| = n$ and a labeling $l : X \to [m]$. The m-uple $x = (x_1, \ldots, x_m) \in X^m$ is *almost completely labeled* if $\{j \in [n] \mid x_i = j \text{ for some } i \in [m]\} = [m] \setminus \{k\}$ for exactly one $k \in [m]$. That is, all labels appear once in x, except for the *missing label k* and a *duplicate label $h \in [m]$* that appears twice.

It's easy to see that if we pivot from an almost completely labeled facet (or vertex) on the duplicate label, or from a completely labeled facet (or vertex) on any label, we reach either an almost completely labeled or a completely labeled facet (or vertex).

We now focus on the case of pivoting on vertices of simple polytopes with labeled facets, that is, the case of the classic Lemke-Howson algorithm, first given by Lemke and Howson in [9]; we follow the very clear exposition given by Shapley in [19].

Running the Lemke-Howson algorithm defines a *Lemke path* that connects two different completely labeled vertices through almost completely labeled vertices and edges where the only missing label is k.

Proposition 2.1. *The Lemke-Howson algorithm 1 returns a solution to the problem* ANOTHER COMPLETELY LABELED VERTEX.

Furthermore, the number of completely labeled vertices in a simple polytope with labeled facets is even.

Proof. The fact that x is completely labeled is trivial; we must show that $x \neq x_0$. At each vertex x' of the Lemke path there are only two edges corresponding to the missing label k, since P is simple; one is the edge that has

---

Algorithm 1: Lemke-Howson algorithm

   input : A simple d-polytope P with n facets. A labeling $l_f : [n] \to [d]$

             of the facets of P. A vertex x_0 of P, completely labeled for l.

   output: A completely labeled vertex $x \neq x_0$ of P.

1 choose any label $k \in [d]$

2 pivot on label k from x_0 to x

3 **while** *x is not completely labeled* **do**

4   pivot on the duplicate label h from x to $x' \neq x_0$

5   set $x_0 = x$, $x = x'$

6 **return** x

---

been traversed to get to x', the other one will be traversed to leave it in the next step. Therefore, there are no "loops" where a vertex is visited more than once; Lemke paths are *simple paths*.

Each Lemke path is uniquely determined by its missing label and its starting point; furthermore, the Lemke path from the endpoint with the same missing label will lead back to the starting point. Since the endpoint and the starting point are different, the Lemke paths must connect an even number of points. $\qquad\square$

For each label $k \in [d]$ chosen in line 1 of Algorithm 1, Lemke paths can be seen as edges of an undirected graph that connect vertices corresponding to the completely labeled vertices of P, with a vertex corresponding to x_0 as standard endpoint. This shows the following proposition.

Proposition 2.2. ANOTHER COMPLETELY LABELED VERTEX *is a problem in* **PPA**.

Applying the parity result in proposition 2.1 to the case of a bimatrix game (not necessarily a unit vector game), and remembering that the point $(\mathbf{0}, \mathbf{0})$

corresponds to an "artificial" equilibrium, we have the following result, due to Lemke and Howson [9].

Theorem 8. (Lemke-Howson [9]) *Every non-degenerate bimatrix game has an odd number of Nash equilibria.*

There are two ways of using the Lemke-Howson algorithm to find a Nash equilibrium of a bimatrix game (A, B).

The first way is to "symmetrize" the game as in proposition 1.2. Let $C = \begin{pmatrix} 0 & A \\ B^\top & 0 \end{pmatrix}$ and let $S = \{z \in \mathbb{R}^{m+n} \mid z \geq \mathbf{0}, \ Cz \leq \mathbf{1}\}$ be the polytope associated to the game $(C, C^\top)$. We can label the $2(m + n)$ inequalities defining the facets of of S as $1, \ldots, m + n, 1, \ldots, m + n$ and apply the Lemke-Howson algorithm starting from the vertex $\mathbf{0}$; this will return a Nash equilibrium (z, z) of C, and the corresponding Nash equilibrium $(x, y) = z$ of (A, B).

We can also follow the "traditional" version of the Lemke-Howson algorithm, alternating moves on the best response polytopes P and Q as defined in 1.4, starting from the couple of vertices $(\mathbf{0}, \mathbf{0})$. Since we move in $\mathbb{R}^m$ and $\mathbb{R}^n$ instead of $\mathbb{R}^{m+n}$, this version is much easier to visualize.

Example 2.1. Consider the 3×3 game (A, B) of example 1.2.

$$A = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}, \qquad B = \begin{pmatrix} 0 & 2 & 4 \\ 3 & 2 & 0 \\ 0 & 2 & 0 \end{pmatrix}.$$

The best response polytopes can be represented as the best response regions (see figure 1.5) extended to the origin $\mathbf{0}$, as in figure 2.1; the label "outside" refers to the "back" of the polytope.

The path starts in $(\mathbf{0}, \mathbf{0})$; we drop the label 2 and we move on the polytope P. The label 6 is duplicate; so we drop the label 6 and we move on the polytope Q; an so on until we reach the point x, that is a Nash equilibrium of (A, B).

The dual version of the Lemke-Howson algorithm 1 and of proposition 2.1 is quite straightforward.

Figure 2.1 Lemke path for missing label 2 on the best response polytopes P and Q of game (1.2).

Proposition 2.3. *The dual Lemke-Howson algorithm 2 returns a solution to the problem* ANOTHER COMPLETELY LABELED FACET.

Furthermore, the number of completely labeled facets in a simplicial polytope with labeled vertices is even.

We also have the analogous of proposition 2.2

Proposition 2.4. ANOTHER COMPLETELY LABELED FACET *is a problem in* **PPA**.

Example 2.2. Consider the octahedron with vertices labeled as in figure 2.2. The facet A is completely labeled; dropping the vertex with label 1 we pivot to the completely labeled facet B.

The Lemke paths for the completely labeled facets of the octahedron are shown in figure 2.3.

By theorems 5 and 6, in the case of unit vector games is enough to apply the Lemke-Howson algorithm to the polytope $P^l = \{x \in \mathbb{R}^m \mid x \geq \mathbf{0}, \ B^\top x \leq \mathbf{1}\}$ in (1.9), or the dual Lemke-Howson algorithm to the polytope $Q = \text{conv}(\{e_1, \ldots, e_m\}) \cup \{c_1, \ldots, c_n\})$ in (1.12). Not only this yield a Nash equilibrium, but no potential solutions are "lost" considering the polytope P^l with m labels instead of the product of polytopes $P \times Q$ with $m + n$ labels, as

---

Algorithm 2: Dual Lemke-Howson algorithm

> **input** : A simplicial m-polytope Q with n vertices. A labeling
>
> $l_v : [n] \to [d]$ of the vertices of P. A vertex F_0 of Q, completely
>
> labeled for l.
>
> **output**: A completely labeled facet $F \neq F_0$ of Q.

1 choose any label $k \in [d]$

2 pivot on label k from F_0 to F

3 while x is not completely labeled **do**

4   pivot on the duplicate label h from F to $F' \neq x_0$

**5**   set $F_0 = x$, $F = F'$

6 return x

---

stated in the following theorem by Savani and von Stengel [18]; an analogous result holds for the dual case.

Theorem 9. *Let (U, B) be a unit vector game, with $U = (e_{l(1)} \cdots e_{l(n)})$ for a labeling $l : [n] \to [m]$; let $P = \{x \in \mathbb{R}^m | x \geq \mathbf{0}, \ B^\top x \leq \mathbf{1}\}$ and $Q = \{y \in \mathbb{R}^n | y \geq \mathbf{0}, \ Ay \leq \mathbf{1}\}$, as in 1.4; and let $P^l = \{x \in \mathbb{R}^m \mid x \geq \mathbf{0}, \ B^\top x \leq \mathbf{1}\}$ as in ??. Then the Lemke path on $P \times Q$ for the missing label k projects to a path on P that is the Lemke path on P^l for missing label k if $k \in [m]$, and for missing label $l(j)$ if $k = m + j$ with $j \in [n]$.*

We finally focus on the case of unit vector games where the simplicial polytope Q is cyclic; that is, the case that we can study from the point of view of Gale strings. Consider $s \in G(m, n)$ with d even as a "loop", and let $s(i) = 1$ for an index $i \in [n]$. Then, by Gale evenness condition, there is an odd run of **1** 's in s either on the left or on the right of position i; let j be the first index after this run. A *pivot on s* is then defined as setting $s(i) = 0$ and $s(j) = 1$. Given a labeling $l_s : [n] \to [m]$, we say that we *pivot on label $l(i)$*, *dropping label $l(i)$* and *picking up label $l(j)$*. The *Lemke-Howson for Gale algorithm* is

Figure 2.2   A pivot on the facets of the octahedron.

Figure 2.3   Right: Lemke paths for label 1. Centre: Lemke paths for labels 2 and
3. Left: all Lemke paths.

defined as follows.

We see the correspondence between the Lemke-Howson and the Lemke-Howson for Gale algorithms in the next example.

Example 2.3. Figure 2.4 shows the cyclic polytope $C_4(6)$ with the labeling given in example 1.14. This corresponds to the labeling $l = 123432$, for which there are four completely labeled Gale strings in $G(4, 6)$: $s_A = \mathbf{1111}00$, $s_B = \mathbf{11}0\mathbf{11}0$, $s_C = \mathbf{1}00\mathbf{111}$ and $s_D = \mathbf{1}0\mathbf{11}0\mathbf{1}$, corresponding to the facets A, B, C and D. Pivoting from facet A dropping label 3 yields facet B; analogously,

---

Algorithm 3: Lemke-Howson for Gale algorithm

---

    input : A labeling $l_s : [n] \to [d]$ such that there is a completely labeled

               Gale string $s_0 \in G(d, n)$.

    output: A completely labeled Gale string $s \in G(d, n)$ such that $s \neq s_0$.

1 choose a label $k \in [d]$

2 pivot on label k from s_0 to s

3 **while** *s is not completely labeled* **do**

4   pivot on the duplicate label h from s to $s' \neq s_0$

5   rename $s_0 = s$, $s = s'$

6 **return** s

---

pivoting from $s_A = \mathbf{1111}00$ dropping label 3 yields $s_B = \mathbf{110110}$.

The analogous of Proposition 2.1 and Proposition 2.3 for the Lemke-Howson for Gale algorithm is the following.

**Proposition 2.5.** *The Lemke-Howson for Gale algorithm 3 returns a solution to the problem* ANOTHER GALE.

Furthermore, the number of completely labeled Gale strings $s \in G(d, n)$, where d is even, is even.

In the case of Gale strings, it is quite easy to orient the Lemke paths, and prove a stronger result of **PPAD** complexity instead of **PPA** as in Proposition 2.2 and Proposition 2.4. We will do so by giving an *index* to the completely labeled Gale strings; we will then show that all the endpoints of the Lemke paths have different index, so the Lemke paths can be oriented from positive to negative index.

LHG oriented

An interesting example, as we will see in the next section, is the following, due to Morris [13].

Example 2.4. Consider the labeling $l = 1234564523$ for $G(6, 10)$ and the com-

Figure 2.4 Lemke-Howson for Gale algorithm: the pivoting to $s_A = \mathbf{111100}$ to $s_B = \mathbf{110110}$ correspond to the pivoting from the facet A to the facet B.

pletely labeled Gale string $s = \mathbf{11111100}$. Dropping the label 1, the Lemke-Howson for Gale algorithm will run as in figure 2.5.

Morris paths are exponentially long

2.2 The Complexity of Gale and Another Gale

We will now give our main result: Another Gale can be solved in polynomial time; therefore, it takes polynomial time to find a Nash Equilibrium of a bimatrix game with dual cyclic best response polytope. Our proof will rely on the construction of a graph and, if possible, a perfect matching for it. A *perfect matching* of a multigraph $G = (V, E)$ is a set $M \subseteq E$ of pairwise non-adjacent edges so that every vertex $v \in V$ is incident to exactly one edge in M. A theorem by Edmonds ([6]) gives the complexity of the associated problem Perfect Matching.

```
       1 2 3 4 5 6 4 5 2 3
      ─────────────────────
       1̲ 1 1 1 1 1 · · · ·
       · 1 1 1̲ 1 1 1̄ · · ·
       · 1 1 · 1̲ 1 1 1̄ · ·
       · 1̲ 1 · · 1 1 1 1̄ ·
       · · 1 1̄ · 1 1̲ 1 1 ·
       · · 1 1 1̄ 1 · 1̲ 1 ·
       · · 1̲ 1 1 1 · · 1 1̄
       · · · 1̲ 1 1 1̄ · 1 1
       · · · · 1̲ 1 1 1̄ 1 1
       1̄ · · · · 1 1 1 1 1
      ─────────────────────
       1 2 3 4 5 6 4 5 2 3
```

Figure 2.5 Morris path on $C(6, 10)$

Perfect Matching

---

input : A multigraph $G = (V, E)$.

output: A perfect matching for G, or No if there is no possible perfect matching for G.

Theorem 10. (Edmonds [6]) *The problem* Perfect Matching *can be solved in polynomial time.*

To prove our main result on Another Gale, we will first focus on the accessory problem Gale, and we will use theorem 10 to prove that it is solvable in polynomial time. We will consider every Gale string as a "loop."

Gale

---

input : A labeling $l : [n] \to [d]$, where d is even and $d < n$.

output: A Gale string $s \in G(d, n)$ that is completely labeled by l

---

49

Theorem 11. *The problem* GALE *is solvable in polynomial time.*

Proof. We give a reduction of GALE to PERFECT MATCHING.

Consider the multigraph $G = (V, E)$ with $V = [d]$, so that the vertices of G correspond to the labels $l(i) \in [d]$, and $E = \{(l(i), l(i + 1))$ for $i \in [n]\}$, so that there is an edge between two vertices if and only if the corresponding labels are next to each other at some index i. Let $s \in G(d, n)$ be a completely labeled Gale string. By Gale evenness condition, every run of s corresponds uniquely to $d/2$ pairs of indices $(i, i + 1)$ with $s(i) = s(i + 1) = 1$, and since s is completely labeled, all labels $l(i) \in [d]$ occur at exactly one of these indices. Then the edges $(l(i), l(i + 1))$ form a perfect matching of G.

Conversely, let $l : [n] \to [d]$ be a labeling, and let M be a perfect matching for G. Consider a bitstring s with $s(i) = s(i + 1)$ for every $(l(i), l(i + 1)) \in M$ and $s(i) = 0$ otherwise. Since M is a matching, all the $(l(i), l(i + 1)) \in M$ are disjoint, so, considering s as a "loop," every run of s is of even length, thus satisfying the Gale evenness condition. Since M is perfect, every vertex $v \in [d]$ is the endpoint of an edge $(l(i), l(i + 1))$, so s has exactly d bits equal to **1**, so it is completely labeled.

We have therefore reduced the problem GALEto PERFECT MATCHING, that by theorem 11 can be solved in polynomial time. $\square$

We give two examples of the construction used in theorem 11.

Example 2.5. Figure 2.6 shows the graph for the Morris labeling $l = 1234564523$, and its two matchings $M = \{e_1, e_3, e_5\}$ and $M' = \{e_8, e_6, e_{10}\}$.

These, in turn, correspond to the completely labeled Gale strings $s = $ **1111110**000 and $s = $ **1**0000**11111**.

A perfect matching for a graph, and therefore a Gale string for a labeling, is not always possible, as shown in the next example.

Example 2.6. Consider the labeling $l = 121314$. The graph G is shown in figure 2.7

**Figure 2.6** The graph G and its matchings for the Morris labeling $l = 1234564523$.

Since there aren't any disjoint edges, it's not possible to find a perfect matching for G. We have already seen in example 1.13 that there isn't any possible completely labeled Gale string for $l = 121314$.

We finally extend the proof of theorem 11 to ANOTHER GALE.

Theorem 12. *The problem* ANOTHER GALE *is solvable in polynomial time.*

Proof. Let $G = (V, E)$ be the graph corresponding to the labeling $l : [n] \to [d]$ as in the proof of theorem 11 and let M be the perfect matching of G corresponding to the completely labeled Gale string $s \in G(d, n)$.

If there are two edges $e, e' \in E$ such that $e \in M$, both e and e' have endpoints $l(i), l(i + 1)$, but $e \neq e'$ (recall that G can be a multigraph), the matching $M' = (M \setminus \{e\}) \cup \{e'\}$ is perfect. The corresponding completely labeled Gale string $s' \in G(d, n)$ satisfies $s' \neq s$, since in s the **1**'s corresponding to the labels $l(i), l(i + 1)$ are in the positions given by the edge e, while in s' they are in the positions given by $e' \neq e$. It takes time $d/2$ to check all edges of M, the time required is still polynomial.

51

Figure 2.7 The graph for the labeling $l = 121314$

We now assume that all the edges in every perfect matching M for G don't have a parallel edge. Since by theorem 2.5 there is an even number of completely labeled Gale strings, the existence of s guarantees the existence of another completely labeled Gale string $s' \neq s$ and the corresponding perfect matching $M' \neq M$. Since $M' \neq M$, there is at least one edge $e' \in M$ such that $e' \notin M'$. Consider the $d/2$ graphs $G_i = (V, E_i)$, where $E_i = E \setminus \{e_i\}$ for $e_i \in M$. Since $V(G) = V(G_i)$ and $E(G_i) \subset E(G)$, every perfect matching for one of these G_i is a perfect matching for G as well. With a brute force approach, we look for a perfect matching in each G_i; this will be M'. Since there are $i \in [d/2]$, the time to find it will be still polynomial. $\qquad\square$

We give two examples of the construction of theorem 12.

Example 2.7. The labeling $l = 123432$ gives the graph G in figure 2.8. Suppose that Edmonds' algorithm returns the matching $M = \{e_1, e_3\}$, associated to the completely labeled Gale string $s = \mathbf{111}00$. The edge e_1 has a parallel edge, e_6; we immediately have a second perfect matching in $M' = \{e_3, e_6\}$, associated to the Gale string $s' = \mathbf{101101}$.

A case without parallel edges in the matching is the Morris graph.

Figure 2.8 The graph for the labeling $l = 123432$.

Example 2.8. Consider the Morris graph of example 2.5; suppose that Edmonds' algorithm returns the perfect matching $M = \{e_1, e_3, e_5\}$, as in figure 2.9 right, corresponding to the completely labeled Gale string $s = \mathbf{1111110}000$. We can then delete the edge e_1 to obtain the graph G_1, as in figure 2.9 left. The graph G_1 has a perfect matching $M' = \{e_6, e_8, e_{10}\}$; this is also a perfect matching of G, corresponding to the string $s' = \mathbf{1}0000\mathbf{11111}$.

Figure 2.9 Left: the Morris graph $G = (V, E)$ with the matching $M = \{e_1, e_3, e_5\}$. Right: the graph $G_1 = (V, E \setminus \{e_1\})$ with the matching $M' = \{e_6, e_8, e_{10}\}$.

Chapter 3

Further results

Acknowledgements

appendix/acknowledgments

Bibliography

[1] A. V. Balthasar (2009). "Geometry and equilibria in bimatrix games." PhD Thesis, London School of Economics and Political Science.

[2] M. M. Casetti (2008). "PPAD Completeness of Equilibrium Computation." MSc Thesis, London School of Economics and Political Science.

[3] M. M. Casetti, J. Merschen, B. von Stengel (2010). "Finding Gale Strings." *Electronic Notes in Discrete Mathematics* 36, pp. 1065–1082.

[4] X. Chen, X. Deng (2006). "Settling the Complexity of 2-Player Nash Equilibrium." *Proc. 47th FOCS*, pp. 261–272.

[5] C. Daskalakis, P. W. Goldberg, C. H. Papadimitriou (2006). "The Complexity of Computing a Nash Equilibrium." *SIAM Journal on Computing*, 39(1), pp. 195–259.

[6] J. Edmonds (1965). "Paths, Trees, and Flowers." *Canad. J. Math.* 17, pp. 449–467.

[7] D. Gale (1963). "Neighborly and Cyclic Polytopes." *Convexity, Proc. Symposia in Pure Math.*, Vol. 7, ed. V. Klee, American Math. Soc., Providence, Rhode Island, pp. 225–232.

[8] D. Gale, H. W. Kuhn, A. W. Tucker (1950). "On Symmetric Games." *Contributions to the Theory of Games* I, eds. H. W. Kuhn and A. W. Tucker, *Annals of Mathematics Studies* 24, Princeton University Press, Princeton, pp. 81–87.

[9] C. E. Lemke, J. T. Howson, Jr. (1964). "Equilibrium Points of Bimatrix Games." *J. Soc. Indust. Appl. Mathematics* 12, pp. 413–423.

[10] A. McLennan, R. Tourky (2010). "Imitation Games and Computation." *Games and Economic Behavior* 70, pp. 4–11.

[11] N. Megiddo, C. H. Papadimitriou (1991). "On Total Functions, Existence Theorems and Computational Complexity." *Theoretical Computer Science* 81, pp. 317–324.

[12] J. Merschen (2012). "Nash Equilibria, Gale Strings, and Perfect Matchings." PhD Thesis, London School of Economics and Political Science.

[13] W. D. Morris Jr. (1994). "Lemke Paths on Simple Polytopes." *Math. Oper. Res.* 19, pp. 780–789.

[14] J. F. Nash (1951). "Noncooperative games." *Annals of Mathematics*, 54, pp. 289–295.

[15] C. H. Papadimitriou (1994). *Computational Complexity.* Addison-Wesley, Reading, MA.

[16] C. H. Papadimitriou (1994). "On the Complexity of the Parity Argument and Other Inefficient Proofs of Existence." *J. Comput. System Sci.* 48, pp. 498–532.

[17] R. Savani, B. von Stengel (2006). "Hard-to-solve Bimatrix Games." *Econometrica* 74, pp. 397–429.

[18] R. Savani, B. von Stengel (2015). "Unit Vector Games." arXiv:1501.02243v1 [cs.GT]

[19] L. S. Shapley (1974). "A Note on the Lemke-Howson Algorithm." *Mathematical Programming Study 1: Pivoting and Extensions*, pp. 175–189

[20] L. Vé, B. von Stengel "Oriented Euler Complexes and Signed Perfect Matchings." arXiv:1210.4694v2 [cs.DM]

[21] J. von Neumann, (1928). "Zur Theorie der Gesellschaftspiele." *Mathematische Annalen*, 100, pp. 295–320.

[22] B. von Stengel (2012). "Completely Labeled Facet is NP-complete." Manuscript, 6 pp.

[23] G. M. Ziegler (1995) *Lectures on Polytopes.* Springer, New York.